

Graph Search

Computing 2 COMP1927 17x1

PROBLEMS ON GRAPHS

- What kinds of problems do we want to solve on/via graphs?
 - Is there a simple path from A to B
 - Is the graph fully-connected?
 - Can we remove an edge and keep it fully connected?
 - Which vertices are reachable from v ? (transitive closure)
 - What is the cheapest cost path from v to w ?
 - Is there a cycle that passes through all V ? (tour)
 - Is there a tree that links all vertices (spanning tree)
 - What is the minimum spanning tree?
 - Can a graph be drawn in a plane with no crossing edges?
 - Are two graphs “equivalent”? (isomorphism)

GRAPH SEARCH

- In this part of the course, we examine algorithms for:
 - Connectivity (simple graphs)
 - path finding (simple/directed graphs)
 - minimum spanning trees (weighted graphs)
 - shortest path (weighted graphs)
- And look at generic algorithms for traversing (searching) graphs that can be used to solve a wide range of graph problems which involves:
 - walking along edges and visiting vertices
 - recording e.g. path taken, vertices visited, etc.
- We begin with one of the simplest graph traversals ...

SIMPLE PATH SEARCH

- Problem: is there a path from vertex v to vertex w ?
- As a function:

```
int isPath(Graph g, Vertex v, Vertex w);
```

- Approach to solving problem:
 - examine vertices adjacent to v
 - if any of them is w , then we are done
 - otherwise check try vertices two edges from v
 - repeat looking further and further from v

SIMPLE PATH SEARCH

- Algorithm for path-finding:


```
// is there a path from V to W?
ToVisit = { V }    // vertices to be checked
Visited = { }      // previously checked vertices
while (still some vertices in ToVisit) {
    X = remove a vertex from ToVisit
    Visited += X
    foreach (Y in adjacentVertices(X)) {
        if (Y == W) return TRUE
        if (Y not in Visited) ToVisit += Y
    }
}
return FALSE
```

SIMPLE PATH SEARCH


- *Graph traversal* also commonly referred to as *Graph search*
- Two different approaches to order of searching: breadth-first search (BFS), depth-first search (DFS)
 - DFS follows one path to completion before considering others
 - DFS uses recursion or a stack, and backtracking
 - BFS "fans-out" from the starting vertex ("spreading" subgraph)
 - BFS maintains a queue of to-be-visited vertices

COMPARISON OF BFS/DFS PATH FINDING

- Is there a path from a to h?



Breadth-first Search




Depth-first Search

DFS vs BFS APPROACHES

- DFS and BFS are closely related.
- Implementation differs only their use of a stack or a queue
 - BFS implemented via a **queue** of to-be-visited vertices
 - DFS implemented via a **stack** of to-be-visited vertices (or recursion)
- Both approaches ignore some edges and avoid cycles by remembering previously visited vertices.

EXERCISE: DFS AND BFS TRAVERSAL

- Show the DFS order we visit to determine $\text{isPath}(a,k)$
- Show the BFS order we visit to determine $\text{isPath}(a,k)$
- Assume neighbours are chosen in alphabetical order



DEPTH FIRST SEARCH

Depth first traversal can be described recursively as

```
depth-first(G, V)
{
    mark V as visited
    for each neighbour W of V {
        if (W already visited) continue
        depth-first(G, W)
    }
}
```

DEPTH FIRST SEARCH


Notes about DFS:

- One approach: consider neighbours in ascending order
- Recursion induces back-tracking
- need a mechanism for "marking" vertices ... in fact, we number them as we visit them
(so that we could later trace path through graph)
- Make use of three global variables:
 - count ... counter to remember how many vertices traversed so far
 - pre[] ... array saying order in which each vertex was visited (pre stands for preorder)
 - st[] ... array storing the predecessor of each vertex

DEPTH FIRST SEARCH TREE


- The edges traversed in a graph walk form a tree
- It corresponds to the call tree of the recursive *dfs* function
- Represents the original graph minus any cycles or alternate paths
- We can use a tree to encode the whole search process
- Each time we visit a vertex we record the previous vertex we came from - if the graph is connected this forms a spanning tree
 - We store this in the *st* array

DEPTH FIRST SEARCH (DFS)



```
// Assume we start with dummy Edge {0,0}
// assume we start with count = 0
// pre[v] = -1 for all v
// st[v] = -1 for all v (stores the predecessor)
// assume adjacency matrix representation
void dfsR (Graph g, Edge e) {
    Vertex i, w = e.w;
    pre[w] = count++;
    st[w] = e.v;
    for (i=0; i < g->V; i++){
        if ((g->edges[w][i] == 1) && (pre[i] == -1))
            dfsR (g, mkEdge(g,w,i));
    }
}
}
```

DEPTH FIRST SEARCH (DFS)




	0	1	2	3	4	5	6	7
pre	0	7	1	4	3	5	2	6
st	0	7	0	4	6	3	2	4

- the edges traversed in the graph walk form a tree
- the tree corresponds to the call tree of the depth first search
- and contents of the *st* array


and to the contents of the *st* array - spanning tree

- pre contains the pre-ordering of the vertices




PROPERTIES OF DFS FORESTS

- If a graph is not connected it will produce a spanning forest
 - If it is connected it will form a spanning tree
- we call an edge connecting a vertex with an ancestor in the DFS tree that is not its parent a **back edge**



EXERCISE: DFS TRAVERSAL

- Which vertices will be visited during $\text{dfs}(g)$:



- How can we ensure that *all* vertices are visited?

GRAPH SEARCH FUNCTION


- The graph may not be connected
 - We need to make sure that we visit every connected component:

```
void dfSearch (Graph g) {  
    int v;  
    count = 0;  
    pre = malloc (sizeof (int) * g->nV);  
    st = malloc (sizeof (int) * g->nV);  
    for (v = 0; v < g->nV; v++) {  
        pre[v] = -1;  
        st[v] = -1;  
    }  
    for (v = 0; v < g->V; v++) {  
        if (pre[v] == -1)  
            dfsR (g, mkEdge (g, v, v));  
    }  
}
```

- The work complexity of the graph search algorithm is $O(V^2)$ for adjacency matrix representation, and $O(V+E)$ for adjacency list representation

EXERCISE: DFS TRAVERSAL

- Trace the execution of $\text{dfs}(g,0)$ on:



- What if we were using DFS to search for a path from 0..5? We would get 0-1-2-3-4-5.
- If we want the shortest (least edges/vertices) path we need to use BFS instead. See later slides for this.


NON-RECURSIVE DEPTH FIRST SEARCH

- We can use a stack instead of recursion:

```
void dfs (Graph g, Edge e) {
    int i;
    Stack s = newStack();
    StackPush (s,e);
    while (!StackIsEmpty(s)) {
        e = StackPop(s);
        if (pre[e.w] == -1) {
            pre[e.w] = count++;
            st[e.w] = e.v;
            for (i = 0; i < g->nV; i++) {
                if ((g->edges[e.w][i] == 1) &&
                    (pre[i] == -1)) {
                    StackPush (s,mkEdge(g,e.w,i));
                }
            }
        }
    }
}
```


EXERCISE: DFS TRAVERSAL

- Show the final state of the pre and st arrays after $\text{dfs}(g,0)$:



DFS ALGORITHMS: CYCLE DETECTION

- **Cycle detection:** does a given graph have any cycles?
 - if and only if the DFS graph has back edges, it contains cycles
 - we can easily detect this in the DFS search:




DFS ALGORITHMS: CYCLE DETECTION

- We are only checking for the existence of cycle, we are not returning it

```
//Return 1 if there is a cycle
int hasCycle (Graph g, Edge e) {
    int i, w = e.w;
    pre[w] = count++;
    st[w] = e.v;
    for (i=0; i < g->V; i++){
        if ((g->edges[w][i] == 1) && (pre[i] == -1)) {
            if(hasCycle (g, mkEdge(g,w,i)))
                return 1;
        } else if( (g->edges[w][i] == 1) && i != e.v){
            //if it is not the predecessor
            return 1;
        }
    }
    return 0;
}
```

DFS ALGORITHMS: CONNECTIVITY


- Each vertex belongs to a connected component
- The function `connectedComponents` sets up the array `cc` to indicate which component contains each vertex



Component 0



Component 1



Component 2

cc

0	0	0	1	1	2	2	2	2
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]

DFS ALGORITHMS

○ Connectivity:

- maintain an extra array `cc` for connected components

```
void connectedComponents (Graph g) {
    int v;
    count = 0;
    ccCount = 0;
    pre = malloc (g->nV *sizeof (int));
    cc = malloc (g->nV *sizeof (int));
    st = malloc (g->nV *sizeof (int));

    for (v = 0; v < g->nV; v++) {
        pre[v] = -1;
        st[v] = -1;
        cc[v] = -1;
    }
    for (v = 0; v < g->V; v++) {
        if (pre[v] == -1) {
            connectedR (g, mkEdge (g, v, v));
            ccCount++;
        }
    }
}
```

```
void connectedR (Graph g, Edge e) {
    int i, w = e.w;
    pre[w] = count++;
    st[w] = e.v;
    cc[w] = ccCount;



    for (i=0; i < g->V; i++){
        if ((g->edges[currV][i] == 1) &&
            (pre[i] == -1)) {
            dfsR (g, mkEdge (g, w, i));
        }
    }
}
```


BREADTH-FIRST SEARCH

- What if we want the shortest path between two vertices?
 - DFS doesn't help us with this problem
- To find the shortest path between v and any vertex w
 - we visit all the vertices adjacent to v (distance 1)
 - then all the vertices adjacent to those we visited in the first step (distance 2)

```
breadth-first(G, V)
{
    add V to queue
    while (queue not empty) {
        V = remove head of queue
        mark V as visited
        foreach neighbour W of V {
            if (W already visited) continue
            add W to tail of queue
        }
    }
}
```

BREADTH-FIRST SEARCH



BREADTH-FIRST SEARCH

- We observed previously that we can simply replace the stack with a queue in the non-recursive implementation to get breadth -first search:

```
void bfs (Graph g, Edge e) {
    int i;
    Queue q = newQueue();
    QueueJoin(q,e);
    while (!QueueIsEmpty(q)) {
        e = QueueLeave(q);
        if(pre[e.w] == -1){
            pre[e.w] = count++;
            st[e.w] = e.v;
            for (i = 0; i < g->nV; i++) {
                if ((g->edges[e.w][i] != 0) &&
                    (pre[i] == -1)) {
                    QueueJoin (q,mkEdge(g,e.w,i));
                }
            }
        }
    }
}
```


IMPROVED BREADTH-FIRST SEARCH

- We can mark them as visited as we put them on the queue since the queue will retain their order. Queue will have at most V entries

```
void bfs (Graph g, Edge e) {
    int i;
    Queue q = newQueue();
    QueueJoin (q,e);
    pre[e.w] = count++;
    st[e.w] = e.v;
    while (!QueueIsEmpty(q)) {
        e = QueueLeave(q);
        for (i = 0; i < g->V; i++) {
            if ((g->edges[e.w][i] != 0) && (pre[i] == -1)) {
                QueueJoin (q,mkEdge(g,e.w,i));
                pre[i] = count++;
                st[i] = e.w;
            }
        }
    }
}
```

EXERCISE: BFS TRAVERSAL

- Show the final state of the pre and st arrays after bfs(g,0):



```
breadth-first(G, V)
{
  add V to queue
  while (queue not empty) {
    V = remove head of queue
    mark V as visited
    foreach neighbour W of V {
      if (W already visited) continue
      add W to tail of queue
    }
  }
}
```

Write code to print out the shortest path from 0 to a given vertex v.

BREADTH-FIRST SEARCH

- For one BFS: $O(V^2)$ for adjacency matrix and $O(V+E)$ for adjacency list
- We can do BFS for every node as root node, and store the resulting spanning trees in a $V \times V$ matrix to store all the shortest paths between any two vertices
- To store and calculate these spanning trees, we need
 - memory proportional to $V * V$
 - time proportional to $V * E$
- Then, we can
 - return path length in constant time
 - path in time proportional to the path length